

```
import numpy as np
import gym
import matplotlib.pyplot as plt

env = gym.make("CartPole-v1")
state = env.reset()
state
# położenie, prędkość, kąt, prędkość kątowa

array([-0.04466579,  0.00737135, -0.03683229,  0.03981359])

import keras
from keras.models import Sequential
from keras.layers import Dense
from collections import deque
import tensorflow as tf
import random
```

Funkcja **Q** aproksymowana przez **sieć neuronową** - na wejściu sieci **tensor o kształcie (1,4)**, na wyjściu sieci **tensor o kształcie (1,2)** zawierający **2 wartości Q** odpowiadające **akcjom w lewo i prawo**.

```
model = Sequential()
model.add(Dense(units = 40, input_dim=4, activation='relu'))
model.add(Dense(units = 40, activation = "relu"))
model.add(Dense(units = 2, activation = "linear"))

opt = tf.keras.optimizers.Adam(learning_rate=0.1)
#opt = keras.optimizers.SGD(learning_rate=0.001)

model.compile(loss='MSE',optimizer=opt)
model.summary()
```

```
Model: "sequential"
```

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 40)	200
dense_1 (Dense)	(None, 40)	1640
dense_2 (Dense)	(None, 2)	82
Total params: 1,922		
Trainable params: 1,922		
Non-trainable params: 0		

### Parametry uczenia:

```
train_episodes = 500
epsilon = 0.3
gamma = 0.99
max_steps = 200
```

```
epsilon = 1
```

### Pętla treningowa:

```
memory = deque(maxlen=100)
```

```
batch_size = 10
```

```
def train():
    state_batch, Qs_target_batch = [], []
    minibatch = random.sample(memory, batch_size)
    for state, action, reward, next_state, done in minibatch:
        if done:
            y = reward
        else:
```

```

    y = reward + gamma*np.max(model.predict(next_state)[0])
    Q_target = model.predict(state)
    Q_target[0][action] = y
    state_batch.append(state)
    Qs_target_batch.append(Q_target)
state_batch = np.array(state_batch).reshape(batch_size,4)
Qs_target_batch = np.array(Qs_target_batch).reshape(batch_size,2)
h=model.fit(state_batch,Qs_target_batch,epochs=1,verbose=0)
loss = h.history['loss'][0]
return loss

```

```

#TODO

```

```

Loss = []
Rewards = []

```

```

for e in range(1, train_episodes+1):
    epsilon = epsilon -(1/train_episodes)
    total_reward = 0
    t = 0

    state = env.reset()
    state = np.reshape(state, [1, 4])

    done = False
    while t < max_steps and done == False:
        Qs = model.predict(state)[0]
        if np.random.rand()<epsilon:
            action = env.action_space.sample()
        else:
            action = np.argmax(Qs)
        next_state, reward, done, _ = env.step(action)
        next_state = np.reshape(next_state, [1, 4])
        total_reward += reward

    if done:
        y = reward
    else:
        y = reward + gamma*np.max(model.predict(next_state)[0])

```

```

Q_target = model.predict(state)
Q_target[0][action] = y

h = model.fit(state,Q_target,epochs=1,verbose=0)

loss = h.history['loss'][0]

state = next_state
t+=1

print(e," R=",total_reward," L=",loss)
Rewards.append(total_reward)
Loss.append(loss)
443 R= 23.0 L= 5.307523480408039e-05
444 R= 10.0 L= 1.1674829636376671e-07
445 R= 15.0 L= 9.391660569235682e-06
446 R= 29.0 L= 8.55073521961458e-05
447 R= 14.0 L= 4.835545155401633e-07
448 R= 35.0 L= 0.000213358347536996
449 R= 42.0 L= 182380.421875
450 R= 19.0 L= 3.953080067731207e-06
451 R= 9.0 L= 0.00011025447020074353
452 R= 10.0 L= 0.0005991252255626023
453 R= 27.0 L= 0.0011107134632766247
454 R= 39.0 L= 48316.47265625
455 R= 19.0 L= 0.0006817791145294905
456 R= 24.0 L= 24.642030715942383
457 R= 32.0 L= 5585.802734375
458 R= 54.0 L= 8181.9443359375
459 R= 12.0 L= 0.00046841567382216454
460 R= 28.0 L= 338.01898193359375
461 R= 22.0 L= 7025.31787109375
462 R= 18.0 L= 1488282.125
463 R= 10.0 L= 1423951.125
464 R= 10.0 L= 666483.3125
465 R= 33.0 L= 8084.64306640625
466 R= 31.0 L= 1357134.5
467 R= 26.0 L= 0.13452787697315216

468 R= 24.0 L= 574424.125
469 R= 30.0 L= 214408.484375
470 R= 9.0 L= 0.10796257108449936

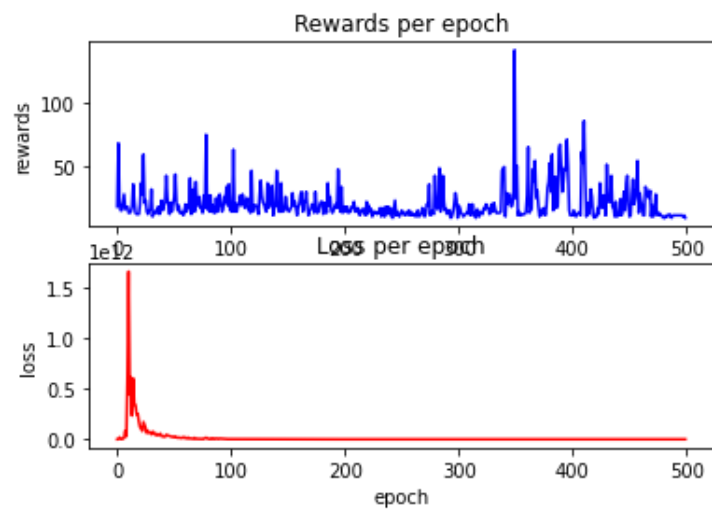
```

```
471 R= 18.0 L= 0.12369771301746368
472 R= 9.0 L= 0.12782898545265198
473 R= 16.0 L= 0.13687244057655334
474 R= 27.0 L= 1183531.75
475 R= 11.0 L= 0.17250114679336548
476 R= 10.0 L= 0.18277494609355927
477 R= 12.0 L= 0.1959000527858734
478 R= 10.0 L= 0.20361244678497314
479 R= 10.0 L= 0.21109023690223694
480 R= 9.0 L= 0.21670474112033844
481 R= 8.0 L= 0.21742995083332062
482 R= 8.0 L= 0.21909686923027039
483 R= 10.0 L= 0.2269432097673416
484 R= 9.0 L= 0.22834771871566772
485 R= 10.0 L= 0.23135773837566376
486 R= 11.0 L= 0.23512154817581177
487 R= 9.0 L= 0.24112603068351746
488 R= 8.0 L= 0.2409575879573822
489 R= 10.0 L= 0.24631565809249878
490 R= 10.0 L= 0.25107094645500183
491 R= 10.0 L= 0.25701117515563965
492 R= 10.0 L= 0.26228341460227966
493 R= 10.0 L= 0.266231894493103
494 R= 10.0 L= 0.26818257570266724
495 R= 10.0 L= 0.27328887581825256
496 R= 10.0 L= 0.27620071172714233
497 R= 10.0 L= 0.2776587903499603
498 R= 9.0 L= 0.27853983640670776
499 R= 10.0 L= 0.27937060594558716
500 R= 8.0 L= 0.27377310395240784
```

```
plt.subplot(211)
plt.ylabel('rewards')
plt.title('Rewards per epoch')
plt.plot(range(len(Rewards)), Rewards, "b")
```

```
plt.subplot(212)
plt.xlabel('epoch')
plt.ylabel('loss')
plt.title('Loss per epoch')
plt.plot(range(len(Loss)), Loss, "r")
```

```
plt.show()
```



✓ 0 s ukończono o 23:42

